

exploration

December 25, 2018

Author: Renjini Ramadasan Nair

Description of the data: Source citation: The data represents the census data from the UCI machine learning database (<http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data>). Detailed information can be read by referring to the original publication (<http://robotics.stanford.edu/~ronnyk/nbtrees.pdf>). There were 15 attributes, of which I have loaded 9 attributes/columns into a dataframe named 'adult' for further analyses. The 9 attributes were the age, workclass, fnlwgt, education-num, sex, capital-gain, capital-loss, hours-per-week and income columns.

```
In [84]: # Import the Numpy and Pandas libraries
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import *
```

```
In [85]: # Download and load data from the url: http://archive.ics.uci.edu/ml/machine-learning
```

```
url = "http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data"
adult1 = pd.read_csv(url, header=None)
```

```
# Add column names
```

```
adult1.columns = ["age", "workclass", "fnlwgt", "education", "education-num", "marital",
```

```
# Select the columns needed to fulfill the current objectives; save as a new dataframe
```

```
adult = adult1.loc[:, ['age', 'workclass', 'fnlwgt', 'education-num', 'sex', "capital
```

0.0.1 Exploration

```
In [86]: adult.head()
```

```
Out[86]:
```

	age	workclass	fnlwgt	education-num	sex	capital-gain	\
0	39	State-gov	77516	13	Male	2174	


1	50	Self-emp-not-inc	83311	13	Male	0
2	38	Private	215646	9	Male	0
3	53	Private	234721	7	Male	0
4	28	Private	338409	13	Female	0

	capital-loss	hours-per-week	income
0	0	40	<=50K
1	0	13	<=50K
2	0	40	<=50K
3	0	40	<=50K
4	0	40	<=50K


```
In [87]: adult.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 9 columns):
age                32561 non-null int64
workclass          32561 non-null object
fnlwgt             32561 non-null int64
education-num      32561 non-null int64
sex                32561 non-null object
capital-gain       32561 non-null int64
capital-loss       32561 non-null int64
hours-per-week     32561 non-null int64
income             32561 non-null object
dtypes: int64(6), object(3)
memory usage: 2.2+ MB
```


```
In [88]: plt.hist(data = adult, x = 'age');
```




```
In [89]: plt.hist(data = adult, x = 'fnlwgt')  
         plt.xscale('log');
```




```
In [90]: plt.hist(data = adult, x = 'education-num');
```




```
In [91]: plt.hist(data = adult, x = 'capital-gain')  
plt.xscale('log');
```




```
In [92]: plt.hist(data = adult, x = 'capital-loss')  
         plt.xscale('log');
```




```
In [93]: plt.hist(data = adult, x = 'hours-per-week');
```




```
In [94]: sb.countplot(data = adult, x = 'workclass', color = sb.color_palette()[0])  
plt.xticks(rotation = 50);
```




```
In [95]: sb.countplot(data = adult, x = 'sex', color = sb.color_palette()[0]);
```




```
In [96]: sb.countplot(data = adult, x = 'income', color = sb.color_palette()[0]);
```




```
In [97]: # Sex related to education
sb.violinplot(data= adult, x = 'sex', y= 'education-num', inner = 'quartile');
```



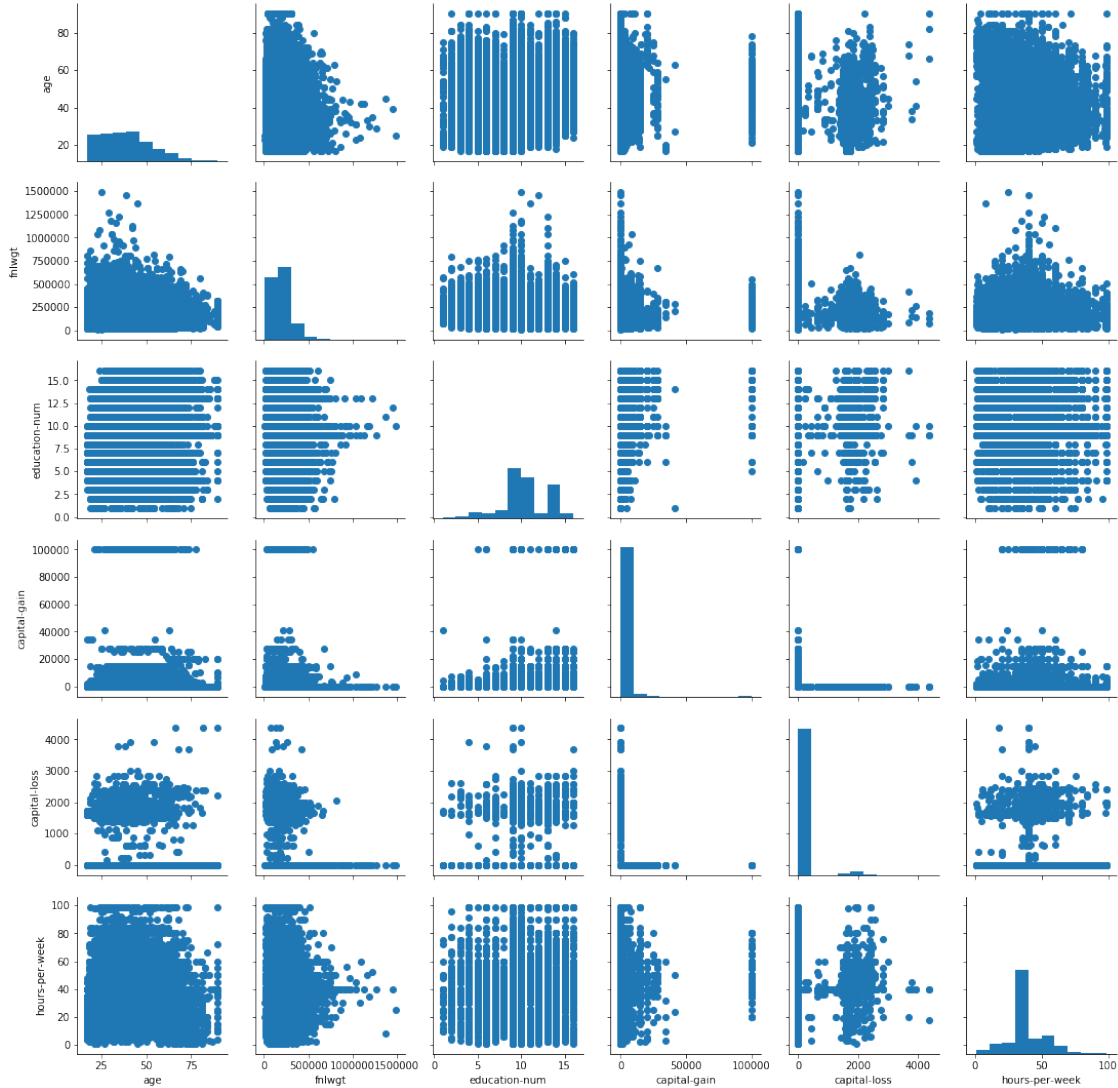
```
In [98]: # Sex related to employment
sb.countplot(data= adult, x = 'workclass', hue= 'sex')
plt.xticks(rotation = 50);
```

```
In [99]: # Age related to employment
sb.boxplot(data= adult, x = 'workclass', y= 'age', color = sb.color_palette()[0])
plt.xticks(rotation = 50);
```



```
In [100]: # Relations between the numerical columns
g = sb.PairGrid(data = adult)
g.map_diag(plt.hist)
g.map_offdiag(plt.scatter);
```



Of the attributes, age, fnlwgt, capital-gain, capital-loss and hours-per-week were numerical attributes and the rest categorical. Initially there were 32561 rows. There were no null values in the numerical data columns, but were present in some categorical columns. Upon exploration it was seen that the education status of males and females followed a similar trend as the employment, being lower for women across job types.

0.0.2 Data Cleaning

I performed data cleaning by following the steps outlined below. The numerical column 'fnlwgt' had outliers, which were replaced with the median value, followed by z-standardization. A plotted histogram showed that this attribute has a more or less normal distribution. For the 'capital-gain' and 'capital-loss' columns, the rows with a few extreme values (>20000 for capital-gain and >3000 for capital loss) were removed. These two columns were also standardized by z-normalization. For the 'education-num', 'workclass' and 'sex' variables, the categories were decoded, imputed and consolidated. The 'education-num' variable was categorized into Primary

(all education lesser than college), Bachelor, Master, Professional and Doctorate. The Doctoral was later consolidated into Professional. The 'workclass' column was categorized into Private, Government and Unemployed categories. The missing values in the 'workclass' column were imputed with the most frequent 'Private' value. The 'sex' column had two categories, 'Male' and 'Female'. Dummy variables were derived for the all three categorical variables and the parental columns deleted. The outliers in the age column were replaced with the median of the non-null values. For the 'age' and 'hours-per-week' columns, the values were binned into 1 (low), 2 (medium) and 3 (high) categories. The obsolete parent columns were deleted.

```
In [101]: # Function definitions
```

```
# Function for checking for outliers in the numeric value columns
```

```
def outlier(var):  
    high = np.mean(var) + 2*np.std(var)  
    low = np.mean(var) - 2*np.std(var)  
    outliers = (var >= high) | (var <= low)  
    return outliers
```

```
# Function to bin numerical columns
```

```
def bins(x, n):  
    BinWidth = (max(x) - min(x))/n  
    bound1 = float('-inf')  
    bound2 = min(x) + 1 * BinWidth  
    bound3 = min(x) + 2 * BinWidth  
    bound4 = float('inf')  
    Binned = np.array([" "] * len(x))  
    Binned[(bound1 < x) & (x <= bound2)] = 1 # Low  
    Binned[(bound2 < x) & (x <= bound3)] = 2 # Med  
    Binned[(bound3 < x) & (x < bound4)] = 3 # High  
    return Binned
```

```
# Function for z-standardization of a numerical column
```

```
def norm(col):  
    x = np.array(col).astype(float)  
    X = pd.DataFrame(x)  
    y = StandardScaler().fit(X).transform(X)  
    return y
```

```
In [102]: # Remove/replace outliers for the numerical columns.
```

```
# Replace outliers with median values for the numerical column, fnlwgt.
```

```
outliers = outlier(adult['fnlwgt'])  
adult.loc[outliers, 'fnlwgt'] = np.median(adult.loc[:, 'fnlwgt'])
```

```
# Remove rows with outlier values; For the capital-gain and capital-loss variables  
# capital-gain
```

```
adult.loc[:, "capital-gain"].sort_values()  
high1 = adult.loc[:, "capital-gain"] > 20000
```

```

adult = adult.loc[~high1, :]
# capital-loss
adult.loc[:, "capital-loss"].sort_values()
high2 = adult.loc[:, "capital-loss"] > 3000
adult = adult.loc[~high2, :]

In [103]: # Normalize values for the numerical columns by z-standardization

# Normalize the fnlwgt column; plot the normalized fnlwgt data.
col = adult["fnlwgt"]
adult["fnlwgt"] = norm(col)

# Normalize the capital-gain column
col = adult["capital-gain"]
adult["capital-gain"] = norm(col)

# Normalize the capital-loss column
col = adult["capital-loss"]
adult["capital-loss"] = norm(col)

In [104]: # Bin the numerical columns (age, hours-per-week) kmeans = KMeans(n_clusters=5).fit(a
# age - Equal-width Binning using numpy
x = np.array(adult['age'])
adult['Binned_age'] = bins(x, 6)
# hours-per-week - Equal-width Binning using numpy
x = np.array(adult['hours-per-week'])
adult['Binned_hours-per-week'] = bins(x, 6)

# Delete obsolete column age and hours-per-week columns
adult.drop('age', axis = 1, inplace = True)
adult.drop('hours-per-week', axis = 1, inplace = True)

In [105]: # Decode education-num column, impute categories and consolidate
# Decode and impute
adult.loc[adult.loc[:, 'education-num'] <=12, 'education-num'] = 'Primary'
adult.loc[adult.loc[:, 'education-num'] ==13, 'education-num'] = 'Bachelor'
adult.loc[adult.loc[:, 'education-num'] ==14, 'education-num'] = 'Master'
adult.loc[adult.loc[:, 'education-num'] ==15, 'education-num'] = 'Professional'
adult.loc[adult.loc[:, 'education-num'] ==16, 'education-num'] = 'Doctorate'
# Consolidate professional and doctorate into professional
adult.loc[adult.loc[:, 'education-num'] =='Doctorate', 'education-num'] = 'Professional'

# Decode the other categorical variable, workclass; impute categories, consolidate a
adult.loc[(adult.loc[:, 'workclass'] == ' State-gov'), 'workclass'] = 'Government'
adult.loc[adult.loc[:, 'workclass'] == ' Self-emp-not-inc', 'workclass'] = 'Unemploy
adult.loc[adult.loc[:, 'workclass'] == ' Federal-gov', 'workclass'] = 'Government'
adult.loc[adult.loc[:, 'workclass'] == ' Local-gov', 'workclass'] = 'Government'
adult.loc[adult.loc[:, 'workclass'] == ' Self-emp-inc', 'workclass'] = 'Private'

```

```

adult.loc[adult.loc[:, 'workclass'] == ' Without-pay', 'workclass'] = 'Unemployed'
adult.loc[adult.loc[:, 'workclass'] == ' Never-worked', 'workclass'] = 'Unemployed'
adult.loc[adult.loc[:, 'workclass'] == ' Private', 'workclass'] = 'Private'
# Private seems highest, so impute '?' with Private
adult.loc[adult.loc[:, 'workclass'] == ' ?', 'workclass'] = 'Private'
# Save the dataframe as a copy for plotting purpose
adult2 = adult.copy()
# add dummy variables
adult.loc[:, "Government"] = (adult.loc[:, "workclass"] == "Government").astype(int)
adult.loc[:, "Private"] = (adult.loc[:, "workclass"] == "Private").astype(int)
adult.loc[:, "Unemployed"] = (adult.loc[:, "workclass"] == "Unemployed").astype(int)
adult.drop('workclass', axis = 1, inplace = True)

# add dummy variables for the categorical column education-num
adult[['Bachelor', 'Master', 'Primary', 'Professional']] = pd.get_dummies(adult['education-num'])
# Delete the obsolete education-num column
adult.drop('education-num', axis = 1, inplace = True)

# add dummy variables for the sex column and delete parent column
adult[['Female', 'Male']] = pd.get_dummies(adult['sex'])
adult.drop('sex', axis = 1, inplace = True)

# Replace the '<=50K' and '>50K' with 1 and 0 respectively, in the income column
adult.loc[(adult.loc[:, 'income'] == ' <=50K'), 'income'] = 1
adult.loc[(adult.loc[:, 'income'] == ' >50K'), 'income'] = 0

```

Following the data cleaning, there were 32297 rows and 15 columns in the 'adult' dataframe. Of the 15, 3 are z-standardized numerical columns named 'fnlwgt', 'capital-gain' and 'capital-loss'. 'Binned_age' and 'Binned_hours-per-week' are binned columns of the original numerical attributes. There were 2, 3 and 4 columns derived by one-hot encoding of the 'sex', 'workclass' and 'education-num' categorical columns respectively. The 'income' column represents the expert labels, classifying each instance into income \leq or $>$ than 50000.